

# External Sorting

## Why can't we simply do quicksort on a disk?

To get  $a[i]$  on

☞ internal memory –  $O(1)$

☞ hard disk



1. find the track;
2. find the sector;
3. find  $a[i]$  and transmit.

*device-dependent*



**Tool: Mergesort**



To simplify –

- ☞ Store data on tapes (can only be accessed sequentially)
- ☞ Can use at least **3** tape drives

[[**Example**]] Suppose that the internal memory can handle  $M = 3$  records at a time.

$T_1$     81 | 94 | 11 | 96 | 12 | 35 | 17 | 99 | 28 | 58 | 41 | 75 | 15

Internal memory    12 | 35 | 96  
*Run*

{  $T_2$     11 | 81 | 94 | 17 | 28 | 99 | 15  
 $T_3$     12 | 35 | 96 | 41 | 58 | 75 |

*Number of passes* = 1+3

$$1 + \lceil \log_2(N / M) \rceil$$

{  $T_1$     11 | 12 | 35 | 81 | 94 | 96 | 15  
 $T_4$     17 | 28 | 41 | 58 | 75 | 99 |

{  $T_2$     11 | 12 | 17 | 28 | 35 | 41 | 58 | 75 | 81 | 94 | 96 | 99 |  
 $T_3$     15

## What are the concerns?

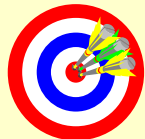
🕒 **Seek** time —  $O(\text{number of passes})$

🕒 Time to **read or write** one **block** of records

🕒 Time to **internally sort**  $M$  records

🕒 Time to **merge**  $N$  records from input buffers to the output buffer

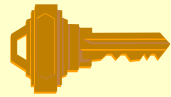
👉 Computer can carry out I/O and CPU processing in **parallel**



## Targets:

- *Reduction of the number of passes*
- *Run merging*
- *Buffer handling for parallel operation*
- *Run generation*

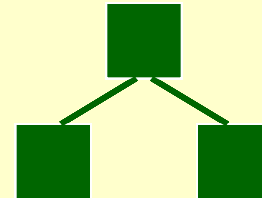
## How to reduce the number of passes?



Use a *k*-way merge!

$T_1$     81 94 11 | 96 12 35 | 17 99 28 | 58 41 75 | 15

$\left\{ \begin{array}{l} T_2 \\ T_3 \\ T_4 \end{array} \right.$ 
 11 81 94 | 41 58 75 |  
 12 35 96 | 15  
 17 28 99 |



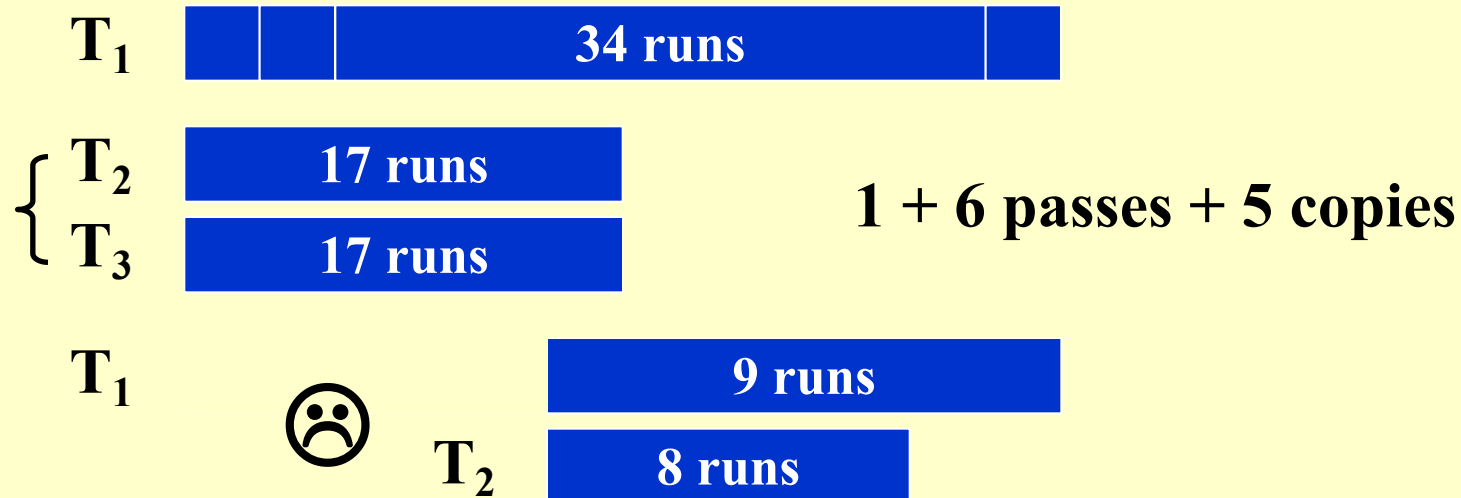
$\left\{ \begin{array}{l} T_1 \\ T_5 \\ T_6 \end{array} \right.$ 
 11 12 17 28 35 81 94 96 99 |  
 15 41 58 75 |

$$\text{Number of passes} = 1 + \lceil \log_k (N / M) \rceil$$

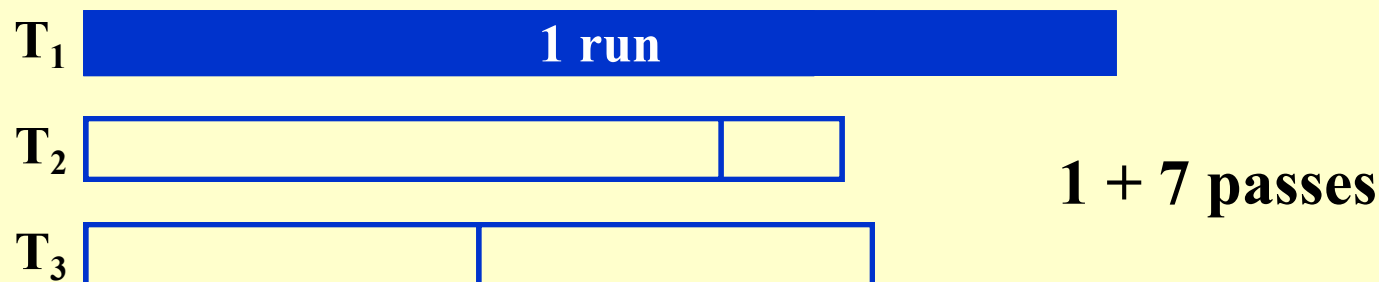


Require *2k* tapes!

Can we use 3 tapes for a 2-way merge?



👉 A smarter way – split *unevenly*



**Discussion 20:**

What will happen if 22 runs are placed on  $T_2$ , with 12 on  $T_3$  ?

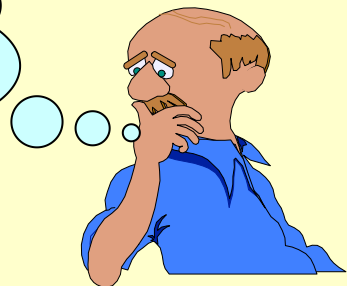
**Claim:** If the number of runs is a Fibonacci number  $F_N$ , then the best way to distribute them is to split them into  $F_{N-1}$  and  $F_{N-2}$  .

**Claim:** For a  $k$ -way merge,  $F_N^{(k)} = F_{N-1}^{(k)} + \dots + F_{N-k}^{(k)}$   
 where  $F_N^{(k)} = 0$  ( $0 \leq N \leq k-2$ ),  $F_{k-1}^{(k)} = 1$

*Polyphase Merge*

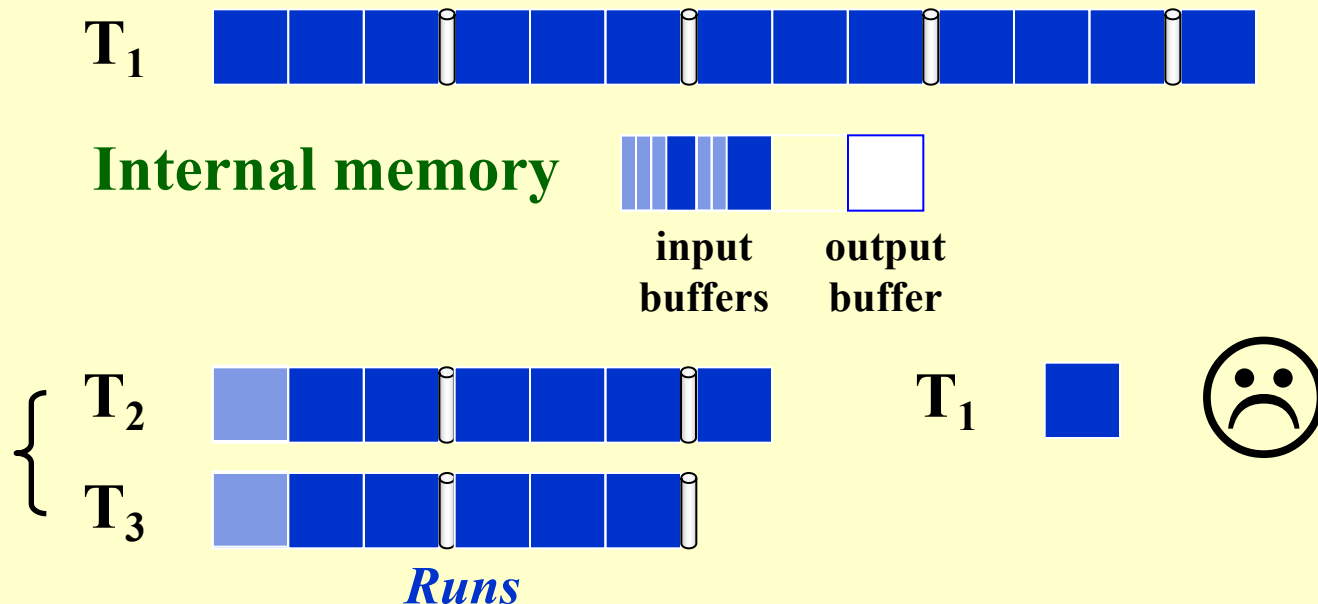
*$k + 1$  tapes only*

What if the initial  
 number of runs is NOT  
 a Fibonacci number?



## How to handle the buffers for parallel operation?

【**Example**】 Sort a file containing **3250** records, using a computer with an internal memory capable of sorting at most **750** records. The input file has a block length of **250** records.

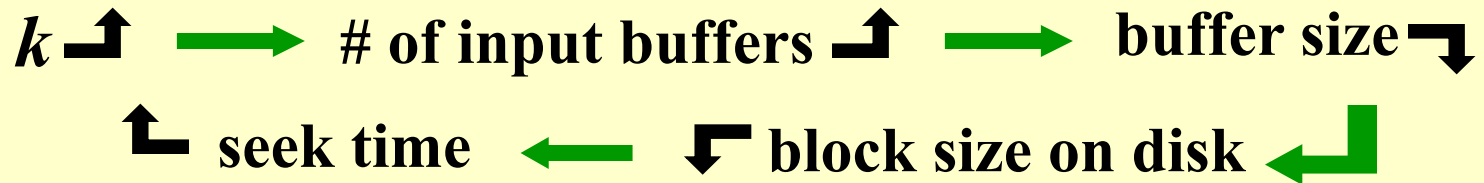




## 2-way merge



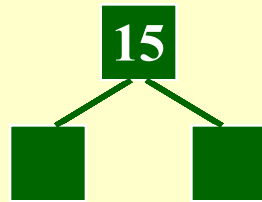
In general, for a  $k$ -way merge we need  $2k$  input buffers and 2 output buffers for parallel operations.



Beyond a certain  $k$  value, the I/O time would actually **increase** despite the decrease in the number of passes being made. The optimal value for  $k$  clearly depends on disk parameters and the amount of internal memory available for buffers.

Can we generate a longer run?

81	94	11	96	12	35	17	99	28	58	41	75	15
----	----	----	----	----	----	----	----	----	----	----	----	----



*Replacement  
selection*

11	81	94	96
----	----	----	----

12	17	28	35	41	58	75	99
----	----	----	----	----	----	----	----

$$L_{avg} = 2M$$

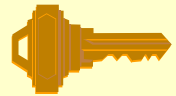
15
----



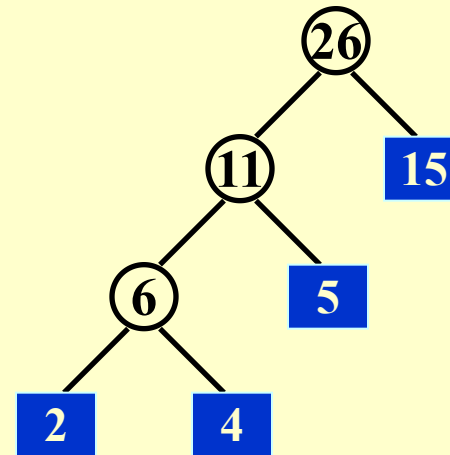
Powerful when input is often *nearly sorted* for external sorting.

## How to minimize the merge time?

【Example】 Suppose we have 4 runs of length 2, 4, 5, and 15, respectively. How can we arrange the merging to obtain **minimum merge times**?



Huffman Tree!



$$2 \times 3 + 4 \times 3 + 5 \times 2 + 15 \times 1 = 43$$

Total merge time =  $O$  ( *the weighted external path length* )

## Reference:

**Data Structure and Algorithm Analysis in C (2<sup>nd</sup> Edition):**  
**Ch.6, p.222-227;** *M.A.Weiss* 著、陈越改编, 人民邮电出版社, 2005

**The Fibonacci Numbers and Polyphase Sorting;** *W. C. LYNCH*, *Case Institute of Technology, Cleveland, Ohio*